

Homework # 02

Lead TAs: Areeba Fatima, Wajid Ali, Sarfraz ,Saif,

Guidelines:

- Attempt all questions by yourself before discussing with peers, this is practice to strengthen your concepts.
- For coding questions, try writing clean, readable code on paper.
- Since this homework is ungraded, focus on learning rather than using LLMs to generate codes and get answers.
- If you get stuck, you are encouraged to:
 - Post your doubts on the course Slack channel.
 - Visit the TAs during office hours for guidance.

Section A: Multiple Choice Questions (MCQs)

Q1) Which traversal of a Binary Search Tree (BST) results in keys in ascending order?

- A) Preorder
- B) Inorder
- C) Postorder
- D) Level order

Q2) A B+ tree supports rapid random access and sequential access because:

- A) Only random access is efficient
- B) Leaves are linked together and tree is balanced
- C) Only sequential access is efficient
- D) It uses hashing for access

Q3) What is the balance factor in an AVL tree?

- A) Height(left subtree) + Height(right subtree)

- B) Height(right subtree) – Height(left subtree)
- C) |Height(left subtree) – Height(right subtree)| ≤ 1
- D) Height(left subtree) – Height(right subtree)

Q4) If you perform a preorder traversal on a BST containing keys {10, 5, 20, 3, 7}, which sequence do you get?

- A) 3, 5, 7, 10, 20
- B) 10, 5, 3, 7, 20
- C) 20, 10, 5, 7, 3
- D) 5, 3, 7, 10, 20

Q5) When inserting a new node into an AVL tree, at most how many rotations are required to restore balance?

- A) O(n)
- B) 2
- C) $\log_2(n)$
- D) Unlimited

Q6) In a B+ tree, where are the actual data records (or pointers to them) stored?

- A) In the internal nodes only
- B) In both internal and leaf nodes
- C) Only in the leaf nodes
- D) Only in the root node

Q7) What is the primary reason for preferring a B+ tree over a BST for indexing a large database stored on disk?

- A) B+ trees are always perfectly balanced.
- B) B+ tree operations have a better Big-O complexity in terms of CPU time.
- C) B+ trees are designed to minimize the number of disk I/O operations.
- D) B+ trees can store more distinct keys than a BST of the same height.

Q8) An AVL tree is a self-balancing version of which fundamental data structure?

- A) B-Tree
- B) Binary Search Tree
- C) Red-Black Tree
- D) Heap

Q9) If a node in an AVL tree has a balance factor of +2, what does this signify according to the formula

- A) The tree is critically left-heavy at that node.
- B) The tree is critically right-heavy at that node.
- C) The tree is perfectly balanced.
- D) The node is a leaf node.

Q10) What is the worst-case time complexity for search, insert, and delete operations in a balanced AVL tree?

- A) $O(n)$
- B) $O(\log n)$
- C) $O(n \log n)$
- D) $O(1)$

Q11) Which of the following is true for a Binary Search Tree (BST)?

- a) The left child is always smaller than the root, and the right child is always larger.
- b) The left and right subtrees are always balanced.
- c) Inorder traversal of a BST gives elements in descending order.
- d) The root must always be the smallest element.

Q12) In a BST, the time complexity of searching for an element in the average case is:

- a) $O(1)$
- b) $O(\log n)$
- c) $O(n)$
- d) $O(n \log n)$

Q13) When deleting a node with two children in a BST, the node is usually replaced with:

- a) Its left child
- b) Its right child
- c) Its inorder successor or predecessor
- d) The root node

Q14) Which traversal of a BST always results in a sorted sequence of elements?

- a) Preorder
- b) Postorder
- c) Inorder
- d) Level-order

Q15) Consider the BST formed by inserting the numbers 15, 10, 20, 8, 12. What is the postorder traversal of this tree?

- a) 8, 12, 10, 20, 15
- b) 12, 8, 10, 20, 15
- c) 8, 10, 12, 20, 15
- d) 20, 12, 10, 8, 15

Q16) The worst-case height of a BST with n nodes is:

- a) $O(\log n)$
- b) $O(n)$
- c) $O(n \log n)$
- d) $O(\sqrt{n})$

Q17) Which of the following sequences cannot represent the inorder traversal of a BST?

- a) 10, 20, 30, 40, 50
- b) 50, 40, 30, 20, 10
- c) 20, 30, 25, 40, 50
- d) 5, 10, 15, 20, 25

Q18) A BST is balanced if:

- a) Every node has at most two children
- b) The left subtree and right subtree of every node differ in height by at most 1
- c) The inorder traversal results in a sorted sequence
- d) The root is the median of all elements

Q19) If you build a BST by inserting 1, 2, 3, 4, 5 in that order, what structure do you get?

- a) A complete binary tree
- b) A skewed tree (right-skewed)
- c) A skewed tree (left-skewed)
- d) A perfectly balanced tree

Q20) If two BSTs contain the same set of elements, then:

- a) Their inorder traversals must be the same
- b) Their structures must be identical
- c) Both a and b
- d) None of the above

Q21) If a BST has height h, the minimum number of nodes it can have is:

- a) h
- b) 2^h
- c) $h + 1$
- d) h

Q22) An AVL tree is a type of:

- a) Binary Search Tree
- b) Heap

- c) Graph
- d) B-tree

Q23) The balance factor of a node in an AVL tree is defined as:

- a) Height of right subtree – Height of left subtree
- b) Height of left subtree – Height of right subtree
- c) Total number of nodes in the tree
- d) Depth of the node

Q24) What are the possible values of the balance factor in a valid AVL tree?

- a) -2, -1, 0, 1, 2
- b) -1, 0, 1
- c) 0 only
- d) Any integer

Q25) The height of an AVL tree with n nodes is always:

- a) $O(n)$
- b) $O(\log n)$
- c) $O(\sqrt{n})$
- d) $O(1)$

Q26) Which of the following is true for an AVL tree?

- a) It is always a complete binary tree
- b) It is always a full binary tree
- c) It is a height-balanced binary search tree
- d) It always has equal number of nodes in left and right subtrees

Q27) Inserting a new node into an AVL tree may require:

- a) At most 1 rotation
- b) At most 2 rotations
- c) At most $\log n$ rotations
- d) Unlimited rotations

Q28) What is the worst-case height H of an AVL tree with n nodes?

- a) $h = O(n)$
- b) $h = O(\log n)$
- c) $h = O(\log \log n)$
- d) $h = O(\sqrt{n})$

Q29) The time complexity of search, insertion, and deletion in an AVL tree is:

- a) $O(1)$
- b) $O(h)$
- c) $O(\log n)$
- d) $O(n)$

Q30) What is the minimum possible height of an AVL tree with 1000 nodes?

- a) ≈ 9
- b) ≈ 10
- c) ≈ 11
- d) ≈ 12

Q31) Which imbalance occurs when a node is inserted into the left subtree of the left child of a node?

- a) LL imbalance
- b) RR imbalance
- c) LR imbalance
- d) RL imbalance

Q32) Which rotation is required to fix an LL imbalance?

- a) Right rotation
- b) Left rotation
- c) Left-Right rotation
- d) Right-Left rotation

Q33) After an insertion, the balance factor of a node becomes +2 and the balance factor of its left child is +1. What rotation is required?

- a) Single left rotation
- b) Single right rotation
- c) Left-Right rotation
- d) Right-Left rotation

Q34) After an insertion, the balance factor of a node becomes +2 and the balance factor of its left child is -1. What rotation is required?

- a) Single left rotation
- b) Single right rotation
- c) Left-Right rotation
- d) Right-Left rotation

Q35) Suppose a B+ Tree of order m has n keys. What is the maximum possible number of leaf nodes?

- a) n
- b) n/m
- c) $n/(m - 1)$
- d) $\log_m(n)$

Q36) In a B+ Tree of order m, the minimum number of keys in a non-root internal node is:

- a) $[m/2] - 1$
- b) $[m/2]$

- c) $m/2$
- d) $[m/2]$

Section B: Short Questions:

Question No.1:

Construct the BST formed by inserting the sequence: 15, 10, 20, 8, 12, 25.

- What is the inorder traversal of the tree?
- What is the height of the BST?

Question No2 :

A node A has balance factor +2. Its left child has balance factor -1.

- a) Which case is this?
- b) Which rotation will fix it?

Question No3:

A B+ Tree of order 4 is used to store integer keys.

- i) What is the minimum and maximum number of keys in a non-root internal node?
- ii) Explain why B+ trees are efficient for **range queries** compared to AVL trees.

Question No4:

Write a C++ function that prints all values in a BST that are strictly greater than a given k.

Question No5:

An e-commerce website wants to implement a feature to filter products by price. Users should be able to see all products within a specific price range (e.g., \$50 to \$200) very quickly. The product database is extremely large and stored on disk.

- a) Which data structure (BST, AVL tree, or B+ tree) would be most suitable for indexing the product prices?
- b) Explain your choice, focusing on the specific requirements of the scenario.
- c) Briefly outline the algorithm to find all products in the price range [minPrice, maxPrice] using your chosen data structure.

Question No6:

You are designing a high-performance in-memory system to check if a username already exists. The system must handle millions of usernames and provide near-instantaneous lookups, insertions (for

new user registrations), and deletions (for account closures). The usernames arrive in a random, unpredictable order.

- a) Which data structure (BST, AVL tree, or B+ tree) would you choose to store the usernames in memory?
- b) Justify your choice, explaining why it is superior to the others for this specific use case.
- c) Write a C++ function `bool doesUserExist(Node* root, string username)` that performs the lookup.

Question No7:

You are developing a leaderboard for a popular online game. The system must support thousands of updates per second as players' scores change. It needs to perform three operations very quickly:

Update/Insert: Add a new player or update an existing player's score.

Lookup: Find a specific player's rank (which is based on their score).

Top-K Query: Display the top 10 players.

All leaderboard data is held in the server's memory for maximum speed.

- a) Which data structure (standard BST, AVL tree, or B+ tree) is most appropriate for this in-memory leaderboard?
- b) Justify your choice by explaining how it efficiently handles the required operations compared to the other two options.
- c) Briefly describe the algorithm to retrieve the top 10 players using your chosen data structure.

Question No8:

A modern operating system needs to index all files on a hard drive by their "last modified" timestamp. This will power a search feature allowing users to quickly find all files modified within a specific date range (e.g., "all files modified between June 1 and June 15, 2025"). The number of files is in the millions, and the index itself is stored on the disk.

- a) Which data structure would you use to build this index?
- b) Explain why your choice is superior for this disk-based, range-query-intensive task.
- c) A node in the tree becomes imbalanced with a balance factor of +2. Its left child has a balance factor of -1. What is this imbalance case called, and what sequence of rotations is needed to fix it?

Question No9:

Part a) Construct the BST by inserting the following elements in order.

[50, 30, 70, 20, 40, 60, 80, 35, 45, 65, 85]

Part b:

- What is the height of this BST? (height counted in terms of edges).
- How many leaf nodes are there? List them.
- How many internal nodes (non-leaf) are there?
- Is the BST complete? Explain briefly.

Part c:

- List the ancestors of node 65.
- List the ancestors of node 20.
- What are the descendants of node 30?
- What is the parent of node 45?
- What is the inorder successor of node 40?
- What is the inorder predecessor of node 60?

Question No10:

a) A binary tree has a root node, where the left subtree contains 500 nodes and the right subtree contains 200 nodes. For preorder, inorder, and postorder traversals, how many nodes will be visited *before* the root is processed?

b) A binary tree T has the following traversals:

- **Preorder:** P Q R S T U
- **Inorder:** R Q S P U T

Find the **postorder traversal** of T

c) The preorder and postorder traversals of a Binary Search Tree (BST) are given below:

- **Preorder:** 60, 35, 20, 45, 40, 75, 70, 90, 85, 95
- **Postorder:** 20, 40, 45, 35, 70, 85, 95, 90, 75, 60

If the root is at level 0, how many nodes are present at **level 2**?

Question No11:

Construct the follow AVL tree after the following operations

- a) Insert([30, 20, 40, 10, 25, 50, 5, 75, 60])
- b) Insert(55)
- c) Delete(20)
- d) Delete (60, 40)

Question No12:

Given a max degree of 5, construct a B+ after the following operations in order:

- a)Insert([7, 10, 1, 23, 5, 15, 17, 9, 11, 39, 35, 8, 40, 25])
- b) Delete(9)
- c)Delete(7)